

Customer Churn Prediction Using Random Forest Model

Customer churn is well known problem in business environment. By customer churn we mean situation when returning clients stop buying company's goods or services. We determine customer churn rate as a percentage of such customers to total number of them. To manage churn rate is important, because for company it costs less to retain customer then to find new one. Another crucial factor to consider is profit. Returning customers usually spend 67% more on company's product and services. If company increases customer retention (or decreases churn rate) by 5%, then it will result in 25% profit increase. If classify customers correctly, then we can take preemptive measures and try to retain customers. Therefore, churn rate prediction plays significant role in firm's profit maximization.

Brief data description and preparation

In this research we use data on clients of telecom company. Data contains characteristics of 7043 customers. Each customer has the following features: ID, gender, seniority status, partner, dependents, tenure term, phone service, multiple lines, internet service, online security, online backup, device protection, technical support, TV streaming, movies streaming, contract term, paperless billing, payment method, monthly charges, total charges, churn. Before proceeding to model's estimation we make data examination. We get rid of missing values: there were 11 observations, which contained null values in any column. We also transform categorical variables into factors, so R built-in models could easily process them. As tenure variable is integer and contains values between 1 and 72 (number of months for which client uses company's services) we decide to split this variable into folds. We split values into 5 folds: less or equal to 1 year, between a year and 2 years, between 2 and 3 years, between 3 and 5 years, more than 5 years. One would expect the more time customer sticks to company less likely it will churn. In this data set 1869 clients terminated their subscriptions, which is approximately 26.58% of total customers number. We provide more detailed data summary in Appendix 2.

Before splitting data into training and test samples we remove unnecessary variables from original dataset. These variables are customer ID (tells us nothing) and tenure (as we created new factor variable). As we get rid of these features out total number of variables becomes 20. We divide our data into training and test sets. 70% of original set is going to be used to train the model. Other 30% we will use to test our model's performance. In these data set churn is our feature we will try to predict. As we have dominance of one class (non-churn – 73.43%) over the other (churn – 26.57%) we should divide our data in training and test sets properly. We assume data set represents true distribution of this variable. Thus, we may want to divide data in such way that we keep the same proportion of churn and non-churn customers in both sets. We implement stratification, so we will not be caught by situation where we have disproportionally large share of churn customers in one set and nearly absence of them in another.

Model's selection rationale

We choose to use random forest model for this prediction problem for several reasons. First, such type of model usually demonstrates good performance in classification problems. Second, random forest model does not depend on variables' distribution. We do not need to satisfy those assumptions. And the most important advantage is the fact that this model does not suffer from overfitting if we set the only one parameter to recommended value (number of randomly selected variables at each split – square root of total number of variables). Random forest model is based on decision trees. Decision tree is series of questions with binary answers (yes / no), which in the end will lead to a certain prediction. We can think

of random forest as ensemble of individual decision trees. Relatively low correlation or its absence between decision trees is key point in random forest model. This gives us another advantage of this model: many uncorrelated individual models gathered as ensemble perform better. By implementing random forest, we effectively build some number of decision trees, which use different samples. Then we split nodes in each tree considering number of randomly selected features. In the last step we obtain ensembled forecast.

Model's specification

For random forest model we must pay attention to one important parameter. We must choose optimal number of randomly selected features at each node split. The idea behind this process is to get least possible node's impurity. This means that the model could definitely classify each observation by using that set of variables. The rule of thumb is to use square root of total number of features for this parameter for classification problem. Random forest function in R sets this parameter's value using that rule by default. We use tuning tool in R package to find optimal value for it (effectively, we cross-validate parameter). By using tuneRF function we find that optimal number of randomly selected variables is 4 (Figure 1). Criteria for optimal value selection is out-of-bag error. We calculate this error by using observations in the training set, which we did not use for model's estimation. This error is just percentage of observations, which model did not predict correctly. We select number of features which leads to minimum value of this error. On Figure 1 we see average out-of-bag error for each number of features as the result of cross-validation. By using rule of thumb, we would get approximately the same result as the total number of features is 19.




Figure 1. Number of randomly selected features optimization

Another parameter of our interest is number of trees. This parameter differs from previous one in the sense that the more trees we grow the better performance model shows. The only cost of increasing this parameter is computational time. This does not lead to any noticeable slowdown in model's estimation for our problem's dimensionality. Thus, we set this parameter to be equal to 1000. These are two parameters, which usually researches pay attention to. But we deal with imbalanced classification problem. Number of churned customers is much less than number of retained customers. For us it is more

important to predict churn class than non-churn. In other words, we consider misclassification of non-churn class less harmful than misclassification of churn class. We can think of it as an over optimistic forecast of company's revenue. When in fact greater number of churned clients will result in less revenue. We should let the model know that despite disproportion between classes, churn class should prevail. We implement this by using "strata" optional argument in random forest function in R. The next step is to identify the proportion by using which model will draw a sample from training set. We choose this to be 45 and 55 respectively. This means that model will draw samples where number of churn class will be greater than non-churn. We estimate such model and obtain the following result:

```
random_forest<-
randomForest(x=x1,y=y1,ntree=1000,mtry=4,sampsize=c(45,55),strata=y1,metric="Accuracy",replace=T)
```

Call:

```
randomForest(x = x1, y = y1, ntree = 1000, mtry = 4, replace = T, strata = y1, sampsize = c(45, 55), metric = "Accuracy")
```

Type of random forest: classification

Number of trees: 1000

No. of variables tried at each split: 4

OOB estimate of error rate: 26.89%

Confusion matrix:	No	Yes	class.error
No	2513	1102	0.3048409
Yes	222	1087	0.1695951

We reach out-of-bag error of 26.86%. For each class errors are 30.48% and 16.96% respectively. Then we run estimated model to obtain prediction for test data and get the following confusion matrix:

testLabel	No	Yes
No	1078	470
Yes	90	470

We can easily derive corresponding errors: 26.57% for out-of-bag error, 30.36% for non-churn class, 16.07% for churn class. As we can see model demonstrates almost the same performance for both training and test sets. For comparison let's see what would have happened if we did not account for disproportion between classes:

```
random_forest1 <- randomForest(x = x1,y = y1, ntree = 1000, mtry = 4,metric="Accuracy",replace=T)
```

Call:

```
randomForest(x = x1, y = y1, ntree = 1000, mtry = 4, replace = T, metric = "Accuracy")
```

Type of random forest: classification

Number of trees: 1000

No. of variables tried at each split: 4

OOB estimate of error rate: 19.88%

Confusion matrix: No Yes class.error

No 3275 340 0.09405256

Yes 639 670 0.48815890

This model has lower OOB error, but the error for churn class is 48.82%. It is nearly 3 times higher than for our selected model (16.07%). By taking into consideration class disproportion we reduced error by 66.67% for churn class and worsen our out-of-bag error only by approximately 25%. For test data it shows almost the same difference in confusion matrix:

testLabel No Yes

No 1391 157

Yes 286 274

After we estimated model and checked its performance on test data, we may be interested in features importance for model (Figure 2):




Figure 2. Feature importance

In the graph above we see features and mean decrease in Gini. We use Gini to measure node impurity:

$$G = \sum_{i=1}^C p(i) * (1 - p(i))$$

where $p(i)$ is probability of picking a sample of class i , C is the total number of classes (2 in our case)

The less this measure the more precisely we can classify our observations. In this graph features are sorted by mean decrease in GINI. We can clearly see 6 most important features: total charges, monthly charges, contract, tenure_folds, payment method, internet service.

In addition, we may want to draw ROC (Receiver Operating Characteristic) curve (Figure 3):

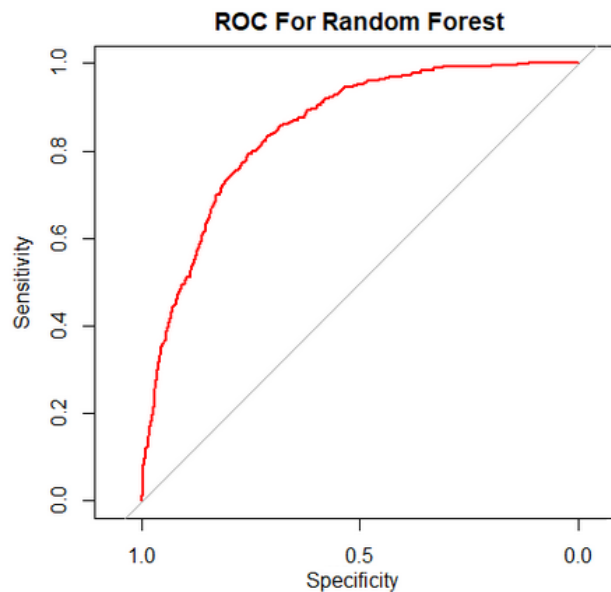


Figure 3. ROC curve

The area under this curve represents model's accuracy. In our case, the accuracy of the model is 84.78%. In fact, we increased model's accuracy by handling imbalanced classification. For plane random forest model accuracy is 82.99%.

Conclusion

In this paper we analyze data on clients of one telecom company. Data includes different clients' features and response variable. Response variable indicates customer's status – churn / non-churn. We choose random forest to predict this response variable on test data. In addition, we noticed that churn class is more important for us to predict accurately. By applying stratified sampling, we improved model's predictive accuracy for positive class (churn). Thus, we achieved classification error for churn class on test data of 16.07%. Genral performance of model is 84.78% (calculated as area under ROC curve).

References

G. James, D. Witten, T. Hastie, R. Tibshirani. "An Introduction to Statistical Learning with Applications in R". Springer Science+Business Media New York, 2013.

Appendix 1 R Script

###installing packages and libraries

install.packages("randomForest")

install.packages("unbalanced")

install.packages("sqldf")

install.packages("ROCR")

install.packages("caret")

install.packages("AUC")

install.packages("mlbench")

install.packages("splitstackshape")

install.packages("mlr")

install.packages("corrplot")

install.packages("gridExtra")

install.packages("ggthemes")

install.packages("party")

install.packages("e1071")

library(pROC)

library(readr)

library(e1071)

library(plyr)

library(corrplot)

library(ggplot2)

library(gridExtra)

library(ggthemes)

library(caret)

library(MASS)

library(randomForest)

library(party)

```

library(unbalanced)

library(sqldf)

library(MASS)

library(ROCR)

library(caret)

library(AUC)

library(mlbench)

library(splitstackshape)

library(mlr)

###load data

data <- read.csv('Customer-Churn.csv')

head(data)

str(data)

### make categorical variables factors if necessary

sapply(data, function(x) sum(is.na(x)))

data <- data[complete.cases(data),]

unique(data['OnlineSecurity'])

f2_col <- c(10:15)

for (i in 1:ncol(data[, f2_col])) {
  data[, f2_col][, i] <- as.factor(mapvalues(data[, f2_col][, i], from = c("No internet service"), to =
c("No")))
}

data$MultipleLines <- as.factor(mapvalues(data$MultipleLines, from = c("No phone service"), to =
c("No")))

data$SeniorCitizen <- as.factor(mapvalues(data$SeniorCitizen, from = c("0", "1"), to = c("No", "Yes")))

```

#may be create even folds with other intervals

```
tenure_folds <- function(tenure){  
  if(tenure >= 0 & tenure <= 12){  
    return('0 - 12 Month')  
  }else if(tenure > 12 & tenure <= 24){  
    return('12 - 24 Month')  
  }else if(tenure > 24 & tenure <= 36){  
    return('24 - 36 Month')  
  }else if(tenure > 36 & tenure <= 50){  
    return('36 - 60 Month')  
  }else if(tenure > 60){  
    return('> 60 Month')  
  }  
}
```

```
data$tenure_folds1 <- sapply(data$tenure, tenure_folds)
```

```
data$tenure_folds1 <- as.factor(data$tenure_folds1)
```

EDA

```
hist(data$MonthlyCharges)
```

```
hist(data$TotalCharges)
```

```
hist(data$tenure)
```

```
summary(data)
```

```
set.seed(100)
```

split dataset

```
train.index <- createDataPartition(data$Churn, p = .7, list = FALSE)
```



```

train <- data[ train.index,]
test  <- data[-train.index,]
trainLabel <- data$Churn[train.index]
testLabel <- data$Churn[-train.index]

summary(train$Churn)
summary(test$Churn)

### drop unnecessary columns
train$customerID <- NULL
train$tenure <- NULL
#train$TotalCharges<-NULL
test$customerID <- NULL
test$tenure <- NULL
#test$TotalCharges<-NULL

str(train)
str(test)

cols <- c(1:18, 20)
x1<-train[,cols]
y1<-train[,19]

### tune parameters on train data set
mtry_opt <- tuneRF(x1, y1, stepFactor=1.5, improve=1e-5, ntree=1000)
print(mtry_opt)

random_forest <- randomForest(x = x1,y = y1, ntree = 1000, mtry = 4,
sampsiz=c(45,55),strata=y1,metric="Accuracy",replace=T)

```

```
random_forest1 <- randomForest(x = x1,y = y1, ntree = 1000, mtry = 4,metric="Accuracy",replace=T)
```

```
### make prediction on test data set
```

```
print(random_forest)
```

```
print(random_forest1)
```

```
plot(random_forest)
```

```
table(testLabel, predict(object = random_forest, newdata = test))
```

```
table(testLabel, predict(object = random_forest1, newdata = test))
```

```
varImpPlot(random_forest)
```

```
### plot ROC
```

```
rf_prediction <- predict(random_forest, test, type="prob")
```

```
ROC_rf <- roc(test$Churn, rf_prediction[,2])
```

```
plot(ROC_rf, col = "red", main = "ROC For Random Forest")
```

```
print(ROC_rf)
```

```
rf_prediction1 <- predict(random_forest1, test, type = "prob")
```

```
ROC_rf1 <- roc(test$Churn, rf_prediction1[,2])
```

```
plot(ROC_rf1, col = "red", main = "ROC For Random Forest")
```

```
print(ROC_rf1)
```

Appendix 2 Data Summary


```
> summary(data)
```

customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0002-ORFBO: 1	Female:3483	No :5890	No :3639	No :4933	Min. : 1.00	No : 680	No :4065
0003-MKNFE: 1	Male :3549	Yes:1142	Yes:3393	Yes:2099	1st Qu.: 9.00	Yes:6352	Yes:2967
0004-TLHLJ: 1					Median :29.00		
0011-IGKFF: 1					Mean :32.42		
0013-EXCHZ: 1					3rd Qu.:55.00		
0013-MHZWF: 1					Max. :72.00		
(other) :7026							


InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract
DSL :2416	No :5017	No :4607	No :4614	No :4992	No :4329	No :4301	Month-to-month:3875
Fiber optic:3096	Yes:2015	Yes:2425	Yes:2418	Yes:2040	Yes:2703	Yes:2731	one year :1472
No :1520							Two year :1685

PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn	tenure_folds1
No :2864	Bank transfer (automatic):1542	Min. : 18.25	Min. : 18.8	No :5163	> 60 Month :1407
Yes:4168	Credit card (automatic) :1521	1st Qu.: 35.59	1st Qu.: 401.4	Yes:1869	0 - 12 Month :2175
	Electronic check :2365	Median : 70.35	Median :1397.5		12 - 24 Month:1856
	Mailed check :1604	Mean : 64.80	Mean :2283.3		24 - 36 Month: 762
		3rd Qu.: 89.86	3rd Qu.:3794.7		36 - 60 Month: 832
		Max. :118.75	Max. :8684.8		

Histogram of data\$MonthlyCharges



Histogram of data\$TotalCharges



Histogram of data\$tenure

